

Groq LPU Acceleration and Sub-Second Inference

Groq LPU Acceleration and Sub-Second Inference

Document ID: doc_22_groq_inference

Topics: groq, inference, latency

1. Overview

This technical specification details groq lpu acceleration and sub-second inference within the MCP PRO orchestration platform.

Detailed sections cover architectural contracts, runtime memory budgets, error handling semantics, and telemetry capture.

2. Implementation Details

The subsystem is encapsulated in the MCP PRO codebase, offering high-throughput, low-latency execution interfaces.

Default parameters are configured for enterprise stability and scalable vector retrieval.

3. Validation and Benchmarking Standards

All performance assertions must be validated through automated benchmark suites recording latency percentiles (mean, p50, p90, p95, p99), error rates, and resource utilization.